

Beyond Relational Model

Monday, January 27, 2025 9:40 AM

Benefits of the Relational Model

- (Mostly) Standardized Data Model and Query Language
- ACID Compliance
 - Atomicity, Consistency, Isolation, Durability
- Works well with highly structured data
- Can handle large amounts of data
- Well understood, lots of tooling, lots of experience

Many way that a RDBMS increases efficiency:

- Indexing
- Directly controlling storage
- Column oriented storage vs row oriented storage
- Query optimization
- Caching/prefetching
- Materializing views
- Precompiled stored procedures
- Data replication and partitioning

Transaction – A sequence of one or more of the CRUD operations performed as a single, logical unit of work

- Either the entire sequence succeeds (COMMIT)
- OR the entire sequence fails (ROLLBACK or ABORT)
- Help ensure...
 - Data integrity
 - Error recovery
 - Concurrency control
 - Reliable data storage
 - Simplified error handling


ACID Properties

- **Atomicity**
 - Transaction is treated as an atomic unit - it is fully executed or no parts of it are executed
- **Consistency**
 - a transaction takes a database from one consistent state to another consistent state
 - consistent state - all data meets integrity constraints
- **Isolation**
 - 2 transactions T1 and T2 are being executed at the same time but cannot affect each other
 - If both T1 and T2 are reading the data – no problem
 - If T1 is reading the same data that T2 may be writing, can result in:
 - Dirty Read
 - Non-repeatable Read
 - Phantom Reads
- **Durability**
 - Once a transaction is completed and committed, its changes are permanent
 - Even in event of system failure, committed transactions are preserved
 - One of the ways this is handled is with logging

Dirty Read

- A transaction T1 is able to read a row that has been modified by another transaction T2 that hasn't yet executed a COMMIT


DIRTY READS



Non-repeatable Read

- 2 queries in a single transaction T1 execute a SELECT but get different values because another transaction T2 has changed data and COMMITTED


NON REPEATABLE READS



Phantom Reads

- when a transaction T1 is running and another transaction T2 adds or deletes rows from the set T1 is using

PHANTOM READS




Flaws with Relational Model

- Sometimes Schemas evolve over time
- Not all apps may need the full strength of ACID compliance
- Joins can be expensive
- A lot of data is semi-structured or unstructured (JSON, XML, etc)
- Horizontal scaling presents challenges
- Some apps need something more performant (real time, low latency systems)

Scalability – Up or Out?

- **Conventional Wisdom:** Scale vertically (up, with bigger more powerful systems) until demands of high availability make it necessary to scale out to some kind of distributed computing model
 - **Distributed system** is "a collection of independent computers that appear to its users as one computer"
 - Characteristics:
 - Computers operate concurrently
 - Computers fail independently
 - No shared global clock
- **But why?** Scaling up is easier - no need to really modify your architecture. But there are practical and financial limits
- **However:** There are modern systems that make horizontal scaling less problematic.

Distributed Data Sources



Replication – making multiple copies

Sharding – splitting on some attribute

- Data is stored on a >1 node, typically replicated
 - So each block of data is available on N nodes
- Can be relational or non-relational
 - MySQL and PostgreSQL support replication and sharding
 - CockroachDB - new player on the scene
 - Many NoSQL systems support one or both models
- Remember: Network partitioning is inevitable!
 - network failures, system failures
 - Overall system needs to be Partition Tolerant
 - System can keep running even w/ network partition


The CAP Theorem

- The CAP Theorem states that it is impossible for a distributed data store to simultaneously provide more than two out of the following three guarantees:
 - Consistency - Every read receives the most recent write or error thrown
 - Availability - Every request receives a (non-error) response - but no guarantee that the response contains the most recent write
 - Can I get to it all the time?
 - Partition Tolerance - The system can continue to operate despite arbitrary network issues.

Consistency

Availability

Partition tolerance



If you cannot limit the number of faults, requests can be directed to any server, and you insist on serving every request, then you cannot possibly be consistent.

- **Consistency + Availability:** System always responds with the latest data and every request gets a response, but may not be able to deal with network issues

ISSUES

- **Consistency + Partition Tolerance:** If system responds with data from a distributed store, it is always the latest, else data request is dropped.
- **Availability + Partition Tolerance:** System always sends response based on distributed store, but may not be the absolute latest data.

